

Ram Srivatsa Kannan

Software Engineer / Systems Researcher



+1 734 741 3242



web.eecs.umich.edu/~ramsri



ramsri@umich.edu




in/ram-srivatsa-kannan-48109a51



ramsrivatsa

Skills

Overview



Programming

Golang

Verilog • VHDL

Python • Bash

C • C++

Hive • SQL

Java

Education

2013 - 2018 **Ph.D. Computer Science**

Ann Arbor, Michigan, USA

Thesis: Enabling Fairness in Cloud Computing Infrastructures

University of Michigan

2009 - 2013 **B.Tech Information Technology**

Madras, Tamil Nadu, India

Thesis: SCOC IP design for Application Specific Architectures

MIT, Anna University

Experience

March 2019 - **Software Engineer**

Uber Technologies

Present

- **Shadow: Safe Rollout System** Designing a distributed system for automated safe roll out of microservices in Uber. Pursuing research on automating integration tests for services as a part of safe rollout pipeline.
- **SubmitQueue: Keeping Master Green at Scale** Optimizing SubmitQueue, a distributed system for maintaining a green masters in a monorepo.

Fall 2013 – Fall 2018 **Graduate Research Assistant**

University of Michigan, Ann Arbor

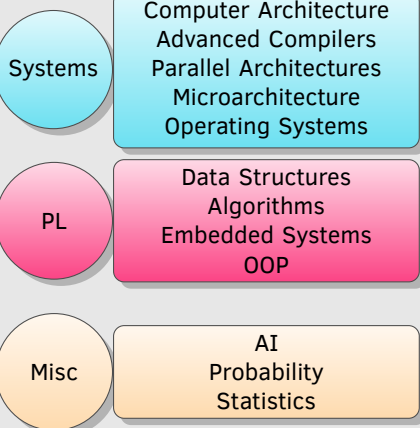
- **Proctor: Identifying Interference in Shared datacenters** - Designed a real-time technique for identifying problematic VMs and its root causes in hyper-converged datacenters that executes CPU, Network and I/O intensive applications.
- **Caliper: Interference Estimator for Multi-tenant Environments Sharing Architectural Resources** - Designed a runtime system to estimate performance degradation of applications running in public clouds and to price users accordingly.
- **GrandSLam – Guaranteeing Datacenter SLAs for Microservices executing in Serverless Computing Infrastructures** - Designed a runtime system to guarantee SLAs among applications utilizing microservers in a serverless computing platform. Implemented DNN based machine learning algorithms in serverless framework using tensorflow on docker containers for this purposes.
- Architecture design for data migration of media files from application to Azure cloud.
- **Big Data analytics on GPUs** - Implemented parallel versions of common data analytics queries and compared it with parallel pthreads versions of the queries in terms of performance. Performed a characterization study of how different input formats of data would affect the performance of queries.
- **Prophet: Precise QoS Prediction on Non-Preemptive Accelerators to Improve Utilization in Warehouse Scale Computers** - Collaborated towards building a GPU based datacenter which guarantees Quality-of-Service for latency-sensitive applications for the state of the art web service applications.
- **Tail-Sniping: Tail Latency Troubleshooting Using Causal Inference** - Collaborated towards building a runtime system to detect performance anomalies using microarchitectural resources and to perform causal analysis with them.

Summer '18 **Research Intern – VMware Research**

VMware Inc.

Worked on a compiler infrastructure that optimizes the compilation times of Verilog in FPGA. Worked on a real time static compiler/scheduler that schedules independent modules present in a Verilog codebase in an optimal order such that emulation time is reduced.

Courses



Interests

Datacenter Design

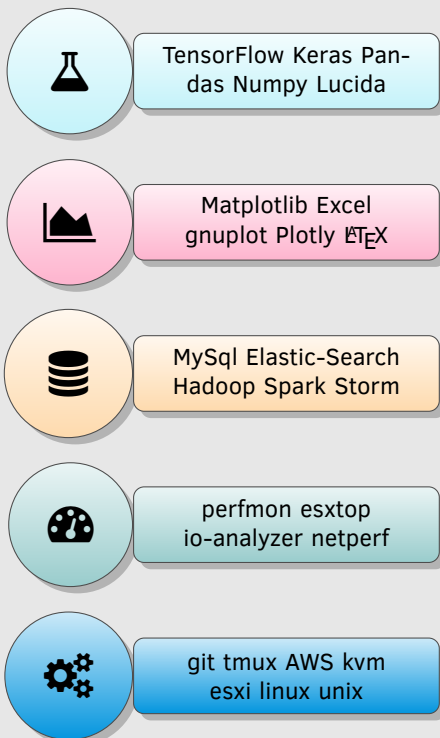
Cloud Computing

Distributes Systems & OS

Software Development

Computer Architecture

Tools



Hobbies



Summer '16

Research Intern – Performance Engineering

VMware Inc.

Worked on real-time, scalable analytics platform which we could automatically detect antagonistic VM behaviors in co-located environments caused due to disk contention using real time system telemetry. Using our technique we were able to identify performance issues and increase individual VM performance by up to $2.5\times$ and $1.8\times$, on an average.

Winter '18

Fall '16, '15

Graduate Student Instructor

University of Michigan

- EECS 483 – Compiler Construction
- EECS 280 – Programming and Data Structures
- EECS 583 – Advanced Compilers

Fall '10 –
Summer '13

Research Trainee

Waran Research Foundation (WARFT), India

- Worked in the high performance computer architecture group that is focused in designing Custom Built heterogeneous Multi-core Architecture - CUBEMACH a node architecture for future generation supercomputing clusters
- Worked on a model for a supercomputing cluster which deals with solving problems in areas like astrophysics, biological sciences, brain modelling and environmental modelling which in future requires high computational power in the range of exa flops.

other projects

Compilers

Loop Invariant Code Motion

Implemented Loop Invariant code motion using an LLVM pass.

Comp Arch

Design and Implementation of Out of order processor

Part of a team which designed and implemented a 2-way superscalar out of order processor in Verilog. It performs out of order execution with Tomasulo's algorithm. It contains additional features like Load Store Queue, tournament branch predictor, return address stack, unified reservation station.

Parallel Comp Arch **Verified MSI and implemented self downgrade optimization**

Verified MSI protocol using a popular tool Murphi. Implemented self downgrade optimization in the existing MSI protocol.

publications

- **Caliper: Interference Estimator for Multi-tenant Environments Sharing Architectural Resource**, Ram Srivatsa Kannan, Michael Laurenzano, Jeongseob Ahn, Jason Mars, Lingjia Tang, ACM Transactions on Architecture and Code Optimization (ACM TACO 2019)
- **GrandSLam - Guaranteeing SLAs for Microservices executing in Serverless Computing Infrastructures**, Ram Srivatsa Kannan, Lavanya Subramanian, Ashwin Raju, Jeongseob Ahn, Jason Mars, Lingjia Tang, European Conference on Computer Systems (Eurosys-19), March 2019
- **Prophet: Precise QoS Prediction on Non-Preemptive Accelerators to Improve Utilization in Warehouse Scale Computers**, Quan Chen, Hailong Yang, Minyi Guo, Ram Srivatsa Kannan, Jason Mars, Lingjia Tang, International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS-17), April 2017
- **Proctor: Identifying Interference in Shared datacenters**, Ram Srivatsa Kannan, Animesh Jain, Michael Laurenzano, Jason Mars, Lingjia Tang, International Symposium on Performance Analysis of Systems and Software (ISPASS-18), April 2018

- **Proctor: Identifying Interference in Shared datacenters**, Ram Srivatsa Kannan, Animesh Jain, Michael Laurenzano, Jason Mars, Lingjia Tang, International Symposium on Performance Analysis of Systems and Software (ISPASS-18), April 2018
- **Compilation Accelerator on Silicon**, Venkateswaran Nagarajan, Ram Srivatsa Kannan et.al., published in the proceedings of 11th IEEE Computer Society Annual Symposium on VLSI, ISVLSI 2012, University of Massachusetts, Amherst
- **SCOC IP Cores for Custom Built Supercomputing Nodes**, Venkateswaran Nagarajan, Ram Srivatsa Kannan et.al., published in the proceedings of 11th IEEE Computer Society Annual Symposium on VLSI, ISVLSI 2012, University of Massachusetts, Amherst

professional activities

External
Reviewer

Conferences

InfoComp 2020, CCGrid 2020, ISCA 2019, PLDI 2018, MICRO 2018, ISCA 2016, HPCA 2016, ASPLOS 2016